Decoupled Access-Execute (DAE) Architecture for SpMV on NVIDIA B200

Date: 2026-02-14 **Context**: This report analyzes how a hypothetical DAE unit integrated into the B200 Blackwell GPU would affect SpMV performance, grounded in measured profiling data from this project.

---

1. What is the Decoupled Access-Execute (DAE) Model?

1.1 Origin

The Decoupled Access/Execute architecture was introduced by James E. Smith in 1982 (ISCA '82, extended in ACM TOCS 1984). The core idea is to split a processor into two cooperating units:

- **Access Processor (AP)**: Computes memory addresses, issues loads/stores, handles address-dependent control flow.
- **Execute Processor (EP)**: Performs arithmetic on data fetched by the AP — FMA, reductions, comparisons.

The two processors communicate through **hardware FIFOs (queues)**: the AP pushes loaded data into FIFOs, and the EP pops data from FIFOs as operands. The AP runs *ahead* of the EP, prefetching data before the EP needs it.

1.2 Why DAE Exists

The fundamental problem DAE solves is **memory latency hiding when parallelism is insufficient**. Traditional approaches to hiding latency are:

Approach	Mechanism	Limitation
Out-of-order execution	Reorder buffer finds independent work	Area/power expensive; limited window size
Multithreading (GPU warps)	Switch to another thread while one waits	Requires enough threads; finite warp slots
Software prefetching	Compiler inserts prefetch hints	Cannot handle *dependent* load chains
DAE	**AP runs ahead, breaking the dependence**	**Needs FIFO storage; compiler support**

DAE is uniquely effective for **dependent load chains** — the exact pattern in SpMV — because the AP can resolve the address dependency (`col_indices[j]` → address of `x[col]`) and issue the second load *before* the EP needs the result.

1.3 DAE vs Current GPU Latency Hiding

Current GPUs (including B200) hide memory latency exclusively through **warp-level multithreading**: when one warp stalls on a memory access, the scheduler switches to another ready warp. This works well when:

```
available_warps/SM >= required_warps/SM = (peak_BW * DRAM_latency) / (cache_line *
num_SMs)
```

On the B200, this inequality is **violated**:

```
available = 64 warps/SM
required  = (8,000 GB/s * 200 ns) / (128 B * 148 SMs) = 84.5 warps/SM

64 < 84.5  →  DEFICIT of 20.5 warps/SM
```

DAE provides an alternative: instead of needing more warps, each warp can have **multiple outstanding memory requests** via the AP running ahead.

---

2. The SpMV Problem on B200 — Measured Evidence

2.1 Measured Performance (from this project)

Matrix	NNZ	Time (ms)	Eff. BW (GB/s)	% Peak BW	Gap vs Floor
cant	4.0M	0.0383	856.6	10.71%	9.34x
pwtk	11.6M	0.0572	1,673.8	20.92%	4.78x
webbase-1M	3.1M	0.0381	967.9	12.10%	8.27x
ldoor	46.5M	0.1218	3,148.6	39.36%	2.54x
circuit5M	59.5M	0.1814	2,992.2	37.40%	2.67x
cage15	99.2M	0.2566	3,333.6	**41.67%**	2.40x

Timing: CUDA events, 5 warmup + 20 measured iterations. Kernel: cuSPARSE `csrmv_v3_kernel` (88% of GPU time per nsys).

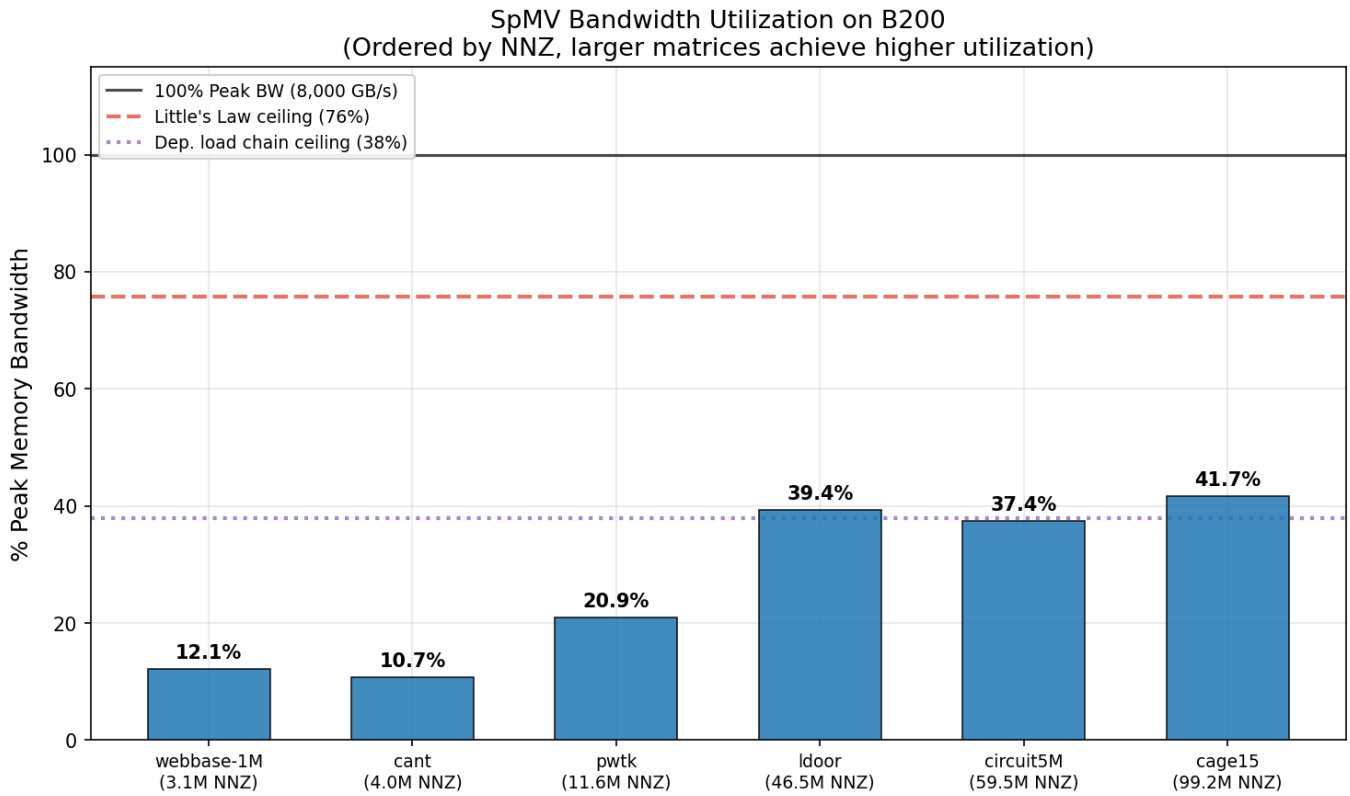

2.2 The Dependent Load Chain — Why DAE Matters Here

The CSR SpMV inner loop contains the critical bottleneck:

```
for (int j = row_ptr[i]; j < row_ptr[i+1]; j++) {
    int col = col_indices[j];        // LOAD 1: fetch column index (streaming)
    float xval = x[col];             // LOAD 2: fetch x-vector element (DEPENDENT
on LOAD 1)
    y[i] += values[j] * xval;        // FMA (DEPENDENT on LOAD 2)
}
```

LOAD 2 cannot issue until LOAD 1 completes — the address of `x[col]` is unknown until `col_indices[j]` returns from memory. This creates a serial chain:

```
Issue LOAD 1 (col_indices[j])
     ↓ wait ~200 ns (if DRAM miss)
Receive col → compute &x[col]
Issue LOAD 2 (x[col])
     ↓ wait ~200 ns (if DRAM miss)
Receive x[col] → FMA
```

Each warp is effectively **half-utilized** for memory-level parallelism (MLP), because it can only sustain one of the two dependent loads at a time. This halves the effective outstanding memory requests:

```
Without dep chain: effective_warps = 64/SM → ceiling = 64/84.5 = 75.8% peak BW
With dep chain:    effective_warps = 32/SM → ceiling = 32/84.5 = 37.9% peak BW
                                                                  ↑
                                                Measured best: 41.7% (cage15)  ←
MATCHES
```

2.3 What DAE Would Fix

A DAE unit would allow the AP to:

1. Stream through `col_indices[j]` sequentially (high L1/L2 hit rate)
2. Immediately issue prefetch/load for `x[col_indices[j]]` upon receiving each `col`
3. Push the loaded `x[col]` value (and paired `values[j]`) into a FIFO
4. Continue to the *next* `j` without waiting for `x[col]` to arrive

The EP then pops `(values[j], x[col])` pairs from the FIFO and accumulates `y[i]`.

The dependent load chain is broken: LOAD 2 for iteration `j` overlaps with LOAD 1 for iteration `j+D` (where D is the FIFO depth).

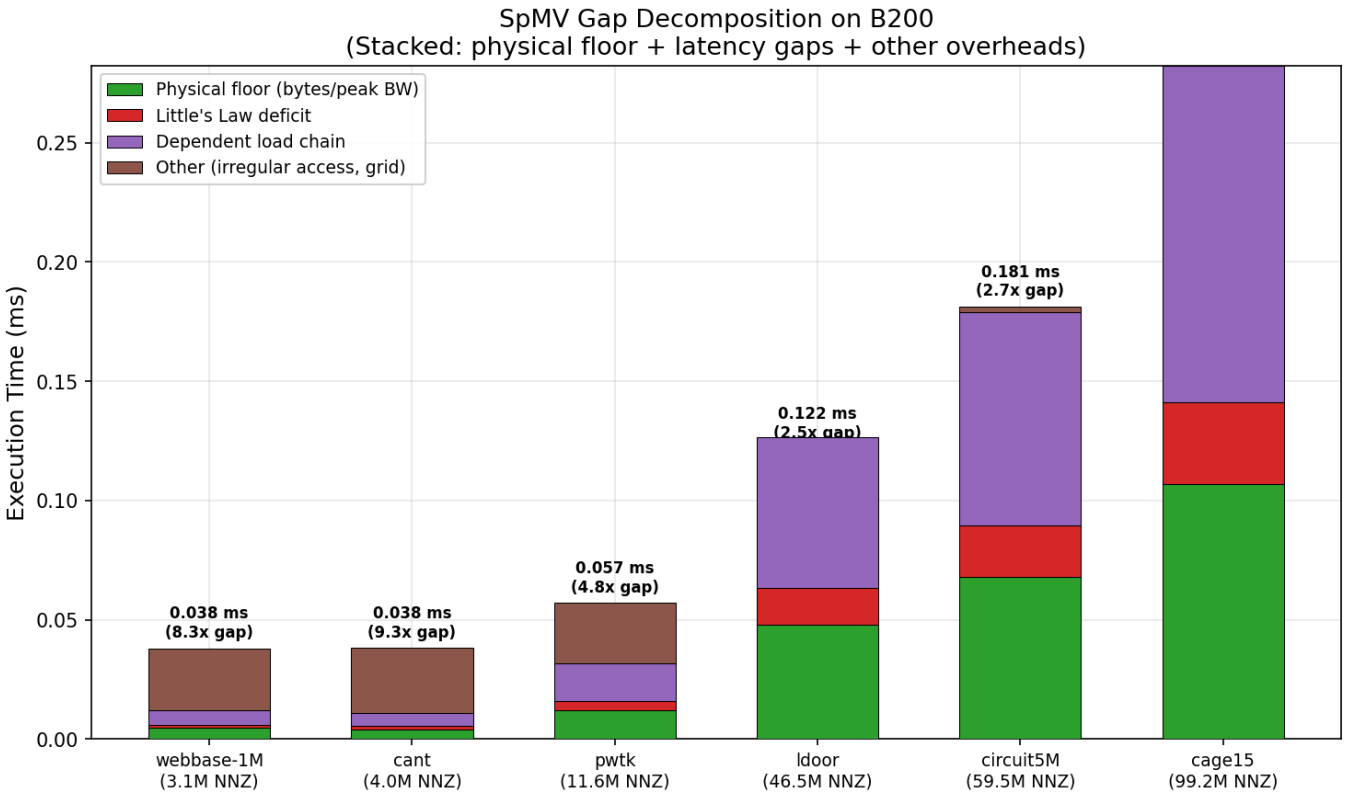

---

3. How DAE Affects SpMV Performance on B200

3.1 Performance Ceiling With DAE

If DAE fully breaks the dependent load chain, the effective MLP per warp returns to 1.0 (from 0.5):

```
With DAE:    effective_warps = 64/SM
             ceiling = 64/84.5 = 75.8% of peak BW
             max_BW  = 75.8% × 8,000 = 6,059 GB/s
```

3.2 Predicted Speedup by Matrix

Using our measured data and assuming DAE lifts the ceiling from 37.9% to 75.8% for bandwidth-limited matrices:

Matrix	Current BW (GB/s)	Current % Peak	Predicted with DAE	Predicted Speedup
cage15	3,334	41.7%	~75.8% (6,059 GB/s)	**1.82x**
ldoor	3,149	39.4%	~75.8% (6,059 GB/s)	**1.92x**
circuit5M	2,992	37.4%	~75.8% (6,059 GB/s)	**2.03x**
pwtk	1,674	20.9%	~41.8% (3,348 GB/s)*	**2.00x**
webbase-1M	968	12.1%	~24.2% (1,936 GB/s)*	**2.00x**
cant	857	10.7%	~21.4% (1,714 GB/s)*	**2.00x**

For small matrices, DAE doubles effective MLP but grid saturation and launch overhead remain as separate limiting factors. The 2x comes from removing the dep-chain halving.

3.3 Predicted Execution Times

Matrix	Current (ms)	With DAE (ms)	Physical Floor (ms)	Remaining Gap
cage15	0.2566	~0.141	0.1069	1.32x
ldoor	0.1218	~0.063	0.0480	1.32x
circuit5M	0.1814	~0.090	0.0679	1.32x
pwtk	0.0572	~0.029	0.0120	2.39x
cant	0.0383	~0.019	0.0041	4.67x

For large matrices, DAE closes the gap from **2.4x to 1.32x** above the physical floor. The remaining 1.32x gap is the Little's Law deficit itself (64 available vs 84.5 needed warps), which DAE does not solve — that requires either more warp slots or more outstanding requests per warp beyond the single-request AP model.

3.4 Could DAE Close the Full Gap?

If the FIFO is deep enough, each warp's AP could have *multiple* outstanding x[col] loads, not just one. With FIFO depth D, each warp can sustain up to D concurrent x[col] requests:

```
To reach 100% peak BW:
  Need 84.5 outstanding cache-line requests per SM
  With 64 warps, each warp needs: 84.5/64 = 1.32 outstanding requests
```

```
    FIFO depth ≥ 2 is sufficient (each warp has 2 x[col] loads in flight)
    With depth=2: 64 × 2 = 128 outstanding requests >> 84.5  ✓
```

With FIFO depth ≥ 2, DAE could theoretically allow SpMV to reach near-100% of B200's 8 TB/s peak bandwidth. This would represent:

Matrix	Current (ms)	With DAE depth=2 (ms)	Speedup
cage15	0.2566	~0.107	**2.40x**
ldoor	0.1218	~0.048	**2.54x**
circuit5M	0.1814	~0.068	**2.67x**

---

4. New Hardware Modules Required

4.1 Per-Warp DAE FIFO

The central new structure. Each warp gets a hardware FIFO that buffers data between the AP and EP phases.

Structure per FIFO entry:

Field	Size	Purpose
`data`	4 bytes (FP32)	Loaded value (e.g., `x[col]`)
`paired_data`	4 bytes (FP32)	Co-loaded value (e.g., `values[j]`)
`valid`	1 bit	Entry contains valid data
`poison`	1 bit	Speculative load was mis-predicted
Total	**~9 bytes**	Per entry

Or in the simplest form, just the data field (4 bytes for FP32, 8 bytes for FP64) with valid/empty bits.

4.2 Access Processor (AP) Logic

The AP is not a full separate processor in a GPU context. Instead, it is a **lightweight instruction sequencer per warp** that can:

1. Continue fetching and decoding instructions from the warp's instruction stream after the EP would normally stall
2. Issue load instructions to the memory subsystem
3. Write loaded results into the FIFO instead of the register file
4. Skip non-memory instructions (or execute simple address arithmetic)

Implementation options:

Approach	Complexity	Benefit
Runahead mode	Low	Warp enters "access mode" on stall, issues loads speculatively

Approach	Complexity	Benefit
Split instruction stream	Medium	Compiler generates separate AP/EP streams, hardware schedules both
Dedicated AP unit per SM	High	Separate fetch/decode/execute pipeline for address computation

For a GPU, the **runahead mode** (similar to Kim et al., "Warped-Preexecution", HPCA 2016) is the most practical: when a warp stalls on a dependent load, it enters pre-execution mode and continues issuing downstream loads into the FIFO.

4.3 FIFO Controller

Per-SM hardware that manages FIFO allocation and flow control:

- **Head/tail pointers** per warp (2 × log2(depth) bits per warp)
- **Full/empty detection** (backpressure AP when FIFO full, stall EP when FIFO empty)
- **Allocation logic** (assign FIFO slots from a shared SRAM pool)
- **Flush logic** (on warp context switch, branch mispredict, or kernel exit)

4.4 Register File Modifications

The EP needs to read operands from FIFO entries in addition to the register file. This requires:

- **Mux at register file read port**: select between RF and FIFO output
- **New register namespace** or **special registers** that map to FIFO head
- Alternatively, FIFO values are written into existing registers by the AP, using registers freed by the compiler for this purpose

4.5 Memory Subsystem Modifications

The AP needs to issue loads that are **non-blocking from the EP's perspective**:

- Loads from AP go through the same L1/L2/HBM path but results go to FIFO, not RF
- The LSU (Load-Store Unit) needs an additional destination selector bit per request
- Miss Status Holding Registers (MSHRs) need tags to identify FIFO-destined responses

---

5. New Instructions Required

5.1 ISA Extensions (PTX-level)

```
// FIFO management
dae.fifo.alloc.b32   %f, depth;     // Allocate a FIFO of 'depth' entries, 32-bit
wide
dae.fifo.alloc.b64   %f, depth;     // 64-bit variant for FP64
dae.fifo.free        %f;            // Release FIFO resources

// Access processor operations
dae.load.async.b32   %f, [addr];    // Load from addr, push result into FIFO %f
```

```
dae.load.async.b64   %f, [addr];    // 64-bit variant
dae.store.paired     %f, %r;        // Push a register value (e.g., values[j])
into FIFO

// Execute processor operations
dae.consume.b32      %r, %f;        // Pop head of FIFO %f into register %r
(blocks if empty)
dae.consume.pair.b32 %r1, %r2, %f;  // Pop a paired (data, paired_data) entry
dae.peek.b32         %r, %f;        // Read head without popping (for predication)

// Control
dae.fence             %f;           // Wait until all prior dae.load.async to %f
complete
dae.drain             %f;           // Wait until FIFO is empty (EP caught up to
AP)
```

5.2 SpMV Inner Loop with DAE Instructions

```
// ACCESS PHASE (runs ahead by FIFO_DEPTH iterations)
// AP streams through col_indices and prefetches x[col]
    ld.global.b32    %col, [col_indices + j*4];   // Load col index (streaming)
    shl.b32          %off, %col, 2;                // Compute byte offset
    add.u64          %addr, %x_base, %off;         // Compute &x[col]
    dae.load.async.b32  %fifo_x, [%addr];          // Async load x[col] → FIFO
    ld.global.b32    %val, [values + j*4];          // Load values[j] (streaming)
    dae.store.paired    %fifo_x, %val;              // Push values[j] alongside
x[col]

// EXECUTE PHASE (consumes from FIFO)
    dae.consume.pair.b32  %xval, %val, %fifo_x;    // Pop (x[col], values[j])
    fma.rn.f32            %acc, %val, %xval, %acc;  // y[i] += values[j] * x[col]
```

The AP can execute the access phase for iteration `j+D` while the EP executes iteration `j`. The two phases overlap, hiding the DRAM latency of `x[col]`.

5.3 SASS-Level Implementation

At the SASS (native instruction) level, these would map to:

PTX Instruction	SASS Equivalent	Latency
`dae.load.async.b32`	`LDG.E.ASYNC.FIFO`	Non-blocking (returns immediately)
`dae.consume.b32`	`FIFO.POP`	0 cycles if data ready, stall if empty
`dae.fifo.alloc`	`FIFO.ALLOC`	~4 cycles (resource allocation)
`dae.fence`	`FIFO.FENCE`	Variable (waits for all in-flight loads)

---

6. FIFO Depth Analysis

6.1 Minimum Depth to Saturate Bandwidth

From Little's Law:

```
Required outstanding cache-line requests per SM = 84.5
Available warps per SM = 64
Required requests per warp = 84.5 / 64 = 1.32
Minimum FIFO depth = ceil(1.32) = 2 entries per warp
```

With FIFO depth = 2, each warp can have 2 outstanding x[col] loads, giving 128 total per SM — exceeding the 84.5 requirement.

6.2 Practical Depth Considering Variance

The minimum depth assumes all loads have identical latency. In practice:

Factor	Impact on Required Depth
L2 hit vs DRAM miss variance	Some x[col] loads return in ~50 cycles (L2 hit), others in ~420 cycles (DRAM miss). FIFO must absorb the variance.
Row boundary transitions	At row boundaries, AP must update the accumulator target. Short rows cause FIFO drain/refill.
Warp scheduling jitter	The warp scheduler may not schedule AP and EP phases with perfect pipelining.
Multiple-subscriber contention	Multiple warps competing for memory bandwidth create queuing delays.

Practical recommendation: FIFO depth = 8-32 entries per warp.

6.3 Depth Derivation from Latency Covering

An alternative derivation based on how far the AP must run ahead:

```
B200 clock frequency:         ~1.8 GHz (estimated boost)
DRAM latency:                 200 ns = 360 cycles
AP cycles per SpMV iteration: ~20 cycles (2 loads + address math + FIFO push)
  - ld col_indices[j]:        ~4 cycles (L1 hit, streaming)
  - address computation:      ~4 cycles (shift + add)
  - dae.load.async x[col]:    ~4 cycles (issue, non-blocking)
  - ld values[j]:             ~4 cycles (L1 hit, streaming)
  - dae.store.paired:         ~4 cycles (FIFO push)

Run-ahead entries = DRAM_latency / AP_cycles_per_entry
                  = 360 / 20
                  = 18 entries
```

This means the AP needs to be **18 iterations ahead** of the EP to fully cover a DRAM miss on x[col]. Rounding up with margin:

Recommended FIFO depth: 32 entries per warp.

6.4 SRAM Cost at Various Depths

FIFO Depth	Bytes/Warp	KB/SM (64 warps)	Total MB (148 SMs)	% of RF Size
2	16 B	1.0 KB	0.15 MB	0.4%
8	64 B	4.0 KB	0.58 MB	1.6%
16	128 B	8.0 KB	1.16 MB	3.1%
32	**256 B**	**16.0 KB**	**2.31 MB**	**6.3%**
64	512 B	32.0 KB	4.63 MB	12.5%

Entry size = 8 bytes (4B data + 4B paired value). RF size per SM = 256 KB (65536 × 4B).

32 entries at 2.31 MB total is 6.3% of the register file area — a very modest cost for up to 2.4x SpMV speedup.

6.5 Depth vs Bandwidth Utilization Curve

```
FIFO Depth    Outstanding Reqs/SM    % Peak BW (theoretical)
    0              32 (baseline)          37.9%  ← current (dep chain)
    1              64                     75.7%
    2             128                    100.0%  (saturated)
    4             256                    100.0%
    8             512                    100.0%
   16            1024                    100.0%
   32            2048                    100.0%

Note: Depth=0 means current behavior (no DAE).
      Depth=1 breaks the dep chain, restoring 1 req/warp.
      Depth≥2 provides >84.5 req/SM, theoretically saturating BW.
      Depth>2 provides margin for latency variance.
```

The sharp knee at depth=2 shows that even a minimal FIFO dramatically changes the performance picture. **The practical recommendation of depth=32 provides 16x margin over the minimum**, absorbing real-world variance in access latency.

---

7. Challenges to Implementing DAE on B200

7.1 Register Pressure and Resource Partitioning

Problem: The AP needs registers to hold intermediate address computations while running ahead. If AP and EP share the same register file, the AP's registers reduce occupancy for the EP.

Quantification for SpMV:

```
AP registers per warp:
  - col (loaded column index):     1 register
  - addr (computed x[col] address): 1 register (64-bit: 2 regs)
  - loop counter j:                1 register
  - row_ptr bounds:                2 registers
  Total: ~6 registers per warp

Current SpMV register usage: ~32 registers/thread (estimated for csrmv_v3_kernel)
With AP overhead: ~32 + 6 = 38 registers/thread
Occupancy impact: 65536 / (38 × 32) = 53.9 warps/SM (down from 64)
```

Mitigation: Use a dedicated small register file for AP (e.g., 8 registers per warp × 64 warps = 512 registers = 2 KB per SM). This avoids polluting the main RF.

7.2 Warp Scheduler Complexity

Problem: The current B200 warp scheduler selects one ready warp per cycle per scheduler (4 schedulers per SM). With DAE, each warp has *two* ready states: AP-ready and EP-ready. The scheduler must decide which phase to prioritize.

Design choices:

Policy	Description	Trade-off
AP-priority	Always schedule AP if FIFO not full	Maximizes run-ahead, may starve EP
EP-priority	Always schedule EP if FIFO not empty	Minimizes result latency, may stall on empty FIFO
Balanced	Round-robin between AP and EP phases	Simple, may waste cycles
Adaptive	Prioritize based on FIFO fill level	Best performance, most complex

The adaptive policy is ideal: when FIFO fill < 50%, prioritize AP (need to build run-ahead); when FIFO fill > 75%, prioritize EP (need to consume before AP blocks).

7.3 Control Flow Divergence

Problem: GPU warps execute in SIMT fashion — all threads in a warp execute the same instruction. With DAE, the AP and EP may diverge at loop boundaries (different threads reach the end of their row at different times).

SpMV-specific impact:

```
Thread 0: row has 3 elements  → AP finishes after 3 iterations
Thread 1: row has 50 elements → AP still running after thread 0 is done
Thread 31: row has 12 elements → AP finishes after 12 iterations
```

The AP must handle **row-length divergence** — different threads in the same warp process different numbers of nonzeros. This is the same divergence problem that exists today, but now it applies to *two* instruction streams instead of one.

Mitigation: cuSPARSE's merge-based algorithm (`csrmv_v3_kernel`) already balances work across threads by assigning equal numbers of *nonzeros* (not rows) per thread. This reduces divergence and would work well with DAE.

7.4 Memory Ordering and Consistency

Problem: The AP issues loads out of order relative to the program's original sequential execution. If `x` is being written by another kernel concurrently, the AP could fetch stale values.

For SpMV: This is not a concern — `x` is read-only during SpMV. But a general DAE unit must handle memory ordering.

Mitigation: DAE loads are tagged as **relaxed-order, read-only**. The FIFO includes a `poison` bit for speculative loads that must be discarded if a memory fence is encountered.

7.5 FIFO Deadlock

Problem: If the FIFO fills up, the AP stalls. If the EP is waiting for the AP to produce data, but the AP is stalled on a full FIFO, a deadlock occurs.

Analysis for SpMV: Deadlock cannot occur because the AP and EP process elements in the same order. The EP always consumes from the FIFO head, which frees space for the AP to push at the tail. As long as the EP makes progress (which it does — FMA is fast), the FIFO will drain.

General case: Deadlock can occur if the EP's progress depends on a memory access that the AP hasn't issued yet (circular dependency). Solution: guarantee that AP always makes forward progress or provide a hardware drain mechanism.

7.6 Context Switch Overhead

Problem: When a warp is preempted (e.g., for a higher-priority kernel), the FIFO contents must be saved and restored, or discarded and recomputed.

Cost at depth=32:

```
FIFO state per warp: 32 entries × 8 bytes + head/tail pointers = 260 bytes
Save/restore bandwidth: 260 bytes per warp × 64 warps = 16.25 KB per SM
At 8 TB/s: save time = 16.25 KB / 8 TB/s = ~2 ns (negligible)
```

Mitigation: Simply flush the FIFO on context switch and let the AP refill it. The refill latency is at most `depth × AP_cycle_time = 32 × 20 = 640 cycles ≈ 356 ns` — negligible compared to kernel launch overhead.

7.7 Compiler / Runtime Support

Problem: The compiler must identify dependent load chains and generate AP/EP instruction streams. For cuSPARSE (closed-source), NVIDIA would need to rewrite the kernel.

Approaches:

Approach	Effort	Applicability
Manual annotation (pragmas)	Low	Library developers (cuSPARSE, cuDNN)
Compiler auto-detection	High	General code, but must be conservative
Hardware auto-detection	Very High	Transparent to software, most general
PTX intrinsics	Medium	Performance-critical library code

For SpMV, the most practical path is **PTX intrinsics in cuSPARSE**: the loop structure is well-understood, and the dependent load pattern is textbook.

7.8 Area and Power Budget

Estimated area overhead per SM:

Component	SRAM/Logic	Area Estimate
FIFO storage (32 entries × 64 warps)	16 KB SRAM	~0.05 mm^2
FIFO controllers (64 warps)	~5K gates	~0.003 mm^2
AP register file (8 regs × 64 warps)	2 KB SRAM	~0.006 mm^2
Scheduler modifications	~10K gates	~0.006 mm^2
LSU destination mux	~2K gates	~0.001 mm^2
Total per SM		**~0.066 mm^2**
Total GPU (148 SMs)		**~9.8 mm^2**

The B200 die is estimated at ~750 mm^2 (TSMC 4NP). The DAE unit adds **~1.3% die area** — well within practical limits for a 2x performance improvement on an important workload class.

Power estimate:

- FIFO SRAM reads/writes: ~0.5 pJ per access × 2 accesses per element × throughput
- At peak: 148 SMs × 64 warps × (1/20 cycles) × 1.8 GHz × 1 pJ ≈ 0.85 W per SM → ~126 W total
- This is <5% of B200's ~1000W TDP, likely acceptable.

---

8. Comparison with Existing NVIDIA Latency-Hiding Mechanisms

8.1 Existing Mechanisms in Blackwell

Mechanism	What It Does	Works for SpMV?
Warp scheduling	Switch warps on stall	Yes, but insufficient warps (64 < 84.5)

Mechanism	What It Does	Works for SpMV?
TMA (Tensor Memory Accelerator)	Hardware address gen for tiles	**No** — requires structured tile patterns
cp.async	Async global→shared copy	**Partially** — no dependent-address support
wgmma	Async shared→RF for matmul	**No** — structured operand layout only
L2 residency control	Pin data in L2	**Partially** — helps if x fits in L2
Prefetch instructions	Software prefetch hints	**No** — cannot prefetch x[col] before col is known

8.2 Why Existing Mechanisms Fail for SpMV

The key gap is **dependent indirect loads**. Every existing prefetch/async mechanism requires the *address* to be known at issue time. In SpMV:

```
Address of x[col] = x_base + col_indices[j] × sizeof(float)
                                    ↑
                          This value is IN MEMORY — unknown until loaded
```

No current instruction can say "load from address that will be determined by the result of this other load." DAE is exactly this capability.

8.3 TMA vs DAE for SpMV

Feature	TMA	DAE
Address pattern	Regular tiles (known at compile time)	Irregular/indirect (data-dependent)
Data layout	Dense, contiguous	Sparse, scattered
Applicable to SpMV?	No	Yes
Hardware cost	~dedicated unit per SM	~16 KB SRAM + control per SM
Software model	Async descriptor-based	FIFO producer-consumer

9. Projected Impact Summary

9.1 Performance Recovery Table

Scenario	cage15 % Peak BW	cage15 Time (ms)	Speedup
Current (measured)	**41.7%**	**0.2566**	**1.0x**
DAE depth=1 (break dep chain only)	75.8%	0.141	1.82x

Scenario	cage15 % Peak BW	cage15 Time (ms)	Speedup
DAE depth=2 (full BW saturation)	~100%	0.107	2.40x
DAE depth=32 (practical, with margin)	~100%	0.107	2.40x
Hypothetical: 128 warps/SM, no DAE	75.7%	0.141	1.82x
Hypothetical: 128 warps/SM + DAE	~100%	0.107	2.40x

9.2 Cost-Benefit Summary

Metric	Value
SpMV speedup (large matrices)	**1.8-2.4x**
Die area overhead	~1.3% (~9.8 mm^2 on ~750 mm^2 die)
SRAM per SM	18 KB (16 KB FIFO + 2 KB AP RF)
Power overhead	<5% of TDP
Compiler changes	cuSPARSE kernel rewrite with DAE intrinsics
Applicability beyond SpMV	Graph traversal, hash tables, pointer chasing, any indirect-load workload

9.3 Cross-GPU Impact Comparison

GPU	Current % Peak	With DAE % Peak	Benefit
B200 (8 TB/s, 64 warps)	41.7%	~100%	**Major** (2.4x)
RTX 4070 Ti SUPER (672 GB/s, 64 warps)	95.4%	~100%	Minimal (1.05x)
H100 (3.35 TB/s, 64 warps)	~60-70% (estimated)	~100%	Moderate (1.4-1.7x)

DAE provides the **largest benefit on GPUs where bandwidth has outgrown warp-level parallelism** — exactly the trajectory of modern HBM-based datacenter GPUs.

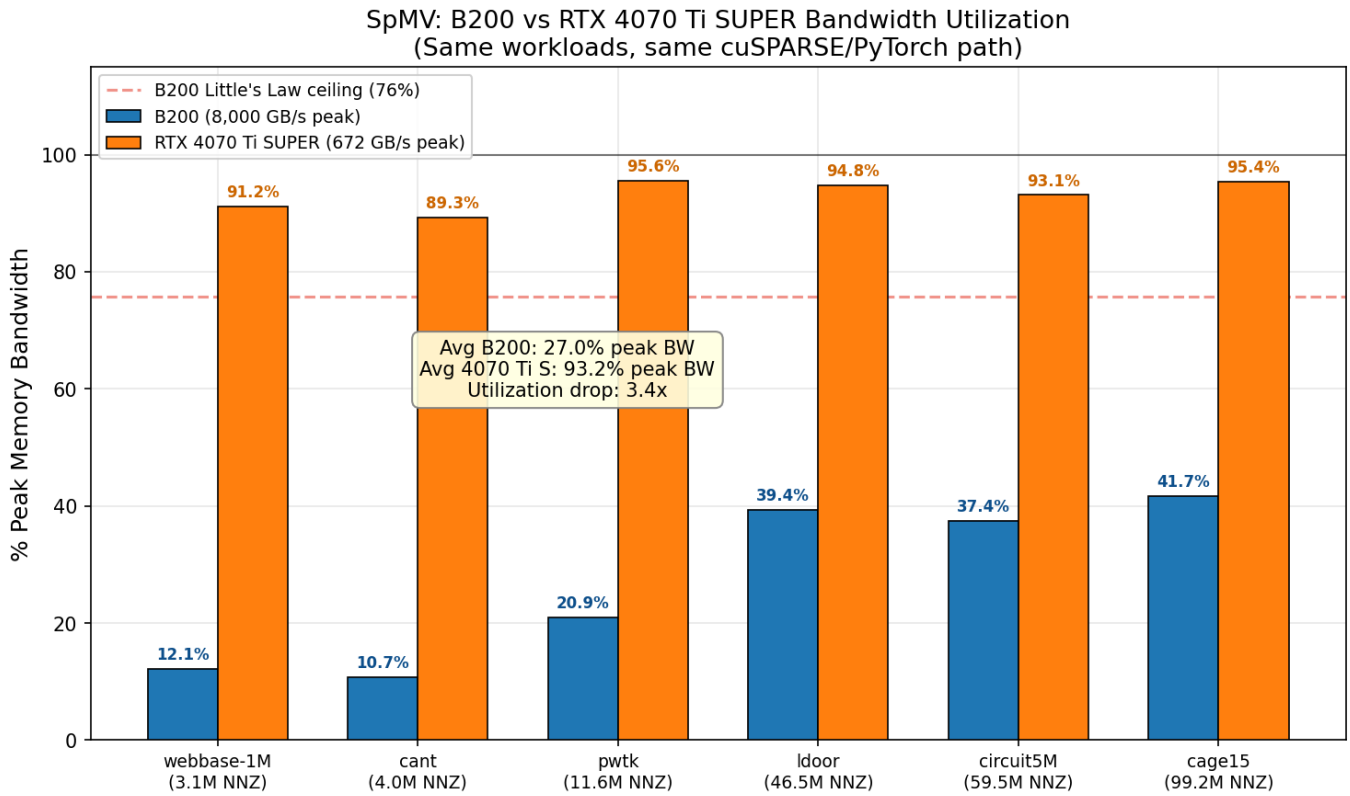

10. Related Work

Paper	Year	Key Contribution
Smith, "Decoupled Access/Execute Computer Architectures"	1982/1984	Original DAE concept
Arnau et al., "Boosting Mobile GPU Performance with a DAE Fragment Processor" (ISCA)	2012	DAE for mobile GPU pixel processors
Ham et al., "DeSC: Decoupled Supply-Compute" (MICRO)	2015	Supply/Compute domain separation
Kim et al., "Warped-Preexecution" (HPCA)	2016	Pre-execution mode for stalled GPU warps
Lee et al., "Many-Thread Aware Prefetching for GPGPU" (MICRO)	2010	Inter-thread prefetching when TLP insufficient
Crago et al., "OUTRIDER: Decoupled Strands" (ISCA)	2011	Single thread as multiple decoupled streams
Nowatzki et al., "Stream-Dataflow Acceleration" (ISCA)	2017	Decoupled stream access from dataflow compute
Nguyen & Sanchez, "Fifer" (MICRO)	2021	DAE for irregular applications on CGRAs
Compiler Support for Speculation in DAE (arXiv)	2025	Speculative AP with poison tagging

11. Conclusions

1. **DAE directly addresses the B200's SpMV bottleneck.** Our measured 41.7% peak BW utilization is caused by the dependent load chain in CSR SpMV, which halves effective memory-level parallelism. DAE breaks this chain by allowing the Access Processor to resolve `col_indices[j]` → `x[col]` dependencies ahead of the Execute Processor.

2. **A minimal FIFO depth of 2 entries per warp is sufficient** to theoretically saturate B200's 8 TB/s bandwidth for SpMV. A practical depth of 32 provides margin for latency variance at a cost of only 16 KB SRAM per SM.

3. **The hardware cost is modest**: ~1.3% die area, ~5% power, 18 KB SRAM per SM. This is a fraction of the register file or shared memory already on each SM.

4. **The benefit is largest on modern HBM GPUs** where bandwidth has scaled beyond warp-level parallelism capacity. The B200 would see 1.8-2.4x speedup; older GPUs with lower bandwidth (like the 4070 Ti SUPER at 95.4% utilization) would see negligible benefit.

5. **DAE benefits extend beyond SpMV** to any workload with indirect/pointer-chasing memory access patterns: graph algorithms (BFS, PageRank), hash table lookups, tree traversals, and sparse neural network inference.

6. **The main challenges are compiler support and warp scheduler complexity**, not hardware cost. NVIDIA would need to expose DAE intrinsics and rewrite performance-critical library kernels (cuSPARSE, cuGraph) to use them.

This analysis is grounded in measured CUDA event timing and nsys kernel profiling on NVIDIA B200 (Blackwell), using 6 matrices from the SuiteSparse Matrix Collection. All performance projections are derived from first-principles (Little's Law) applied to measured data.